

Appendix B Model Kernel Bf16 (Page 1/4)

| GPU | Model | M | cuBLASLt Latency (μ s) | cuSPARSELt Speedup Ratio | | | | | | | | |
|------|-------------|-------------|--------------------------------|--------------------------|------|------|------|------|------|------|-------------|------|
| | | | | 2:4 | 2:6 | 2:8 | 2:10 | 2:12 | 2:14 | 2:16 | 2: ∞ | |
| A100 | Llama3.2-1B | 64 | 1.19e+02 | 1.25 | 0.98 | 0.92 | 0.88 | 0.84 | 0.86 | 0.85 | 0.80 | |
| | | 128 | 1.41e+02 | 1.33 | 1.08 | 1.01 | 0.97 | 0.94 | 0.94 | 0.92 | 0.85 | |
| | | 256 | 1.88e+02 | 1.22 | 0.95 | 0.89 | 0.84 | 0.80 | 0.81 | 0.77 | 0.70 | |
| | | 512 | 3.20e+02 | 1.40 | 1.09 | 1.01 | 0.93 | 0.89 | 0.88 | 0.88 | 0.78 | |
| | | 1024 | 5.90e+02 | 1.46 | 1.10 | 1.03 | 0.94 | 0.89 | 0.90 | 0.89 | 0.79 | |
| | | 2048 | 1.14e+03 | 1.39 | 1.02 | 0.95 | 0.87 | 0.82 | 0.81 | 0.80 | 0.71 | |
| | | 4096 | 2.27e+03 | 1.48 | 1.05 | 0.93 | 0.85 | 0.82 | 0.80 | 0.79 | 0.71 | |
| | | 8192 | 4.42e+03 | 1.46 | 1.05 | 0.94 | 0.86 | 0.82 | 0.82 | 0.79 | 0.71 | |
| | | 16384 | 8.85e+03 | 1.43 | 1.03 | 0.93 | 0.86 | 0.81 | 0.82 | 0.79 | 0.70 | |
| | | BitNet-2B | 64 | 1.28e+02 | 1.34 | 1.10 | 1.03 | 1.00 | 0.96 | 0.95 | 0.94 | 0.85 |
| | | | 128 | 1.50e+02 | 1.35 | 1.07 | 1.00 | 0.95 | 0.91 | 0.90 | 0.90 | 0.80 |
| | | | 256 | 1.95e+02 | 1.37 | 1.06 | 0.98 | 0.94 | 0.90 | 0.87 | 0.86 | 0.77 |
| | | | 512 | 3.59e+02 | 1.50 | 1.16 | 1.07 | 1.02 | 0.95 | 0.93 | 0.93 | 0.82 |
| | | | 1024 | 6.52e+02 | 1.52 | 1.17 | 1.05 | 0.98 | 0.91 | 0.89 | 0.90 | 0.77 |
| | | | 2048 | 1.26e+03 | 1.50 | 1.12 | 0.98 | 0.94 | 0.86 | 0.86 | 0.83 | 0.72 |
| | | | 4096 | 2.47e+03 | 1.52 | 1.14 | 0.98 | 0.94 | 0.90 | 0.86 | 0.86 | 0.75 |
| | | | 8192 | 4.93e+03 | 1.50 | 1.10 | 0.97 | 0.91 | 0.87 | 0.84 | 0.83 | 0.74 |
| | | Llama3.2-3B | 64 | 1.73e+02 | 1.35 | 1.14 | 1.01 | 0.98 | 0.93 | 0.90 | 0.88 | 0.79 |
| | | | 128 | 2.02e+02 | 1.44 | 1.16 | 1.03 | 1.00 | 0.95 | 0.92 | 0.92 | 0.80 |
| | | | 256 | 2.92e+02 | 1.34 | 1.03 | 0.91 | 0.87 | 0.84 | 0.81 | 0.79 | 0.70 |
| | | | 512 | 5.01e+02 | 1.49 | 1.15 | 1.02 | 0.98 | 0.93 | 0.90 | 0.88 | 0.76 |
| | | | 1024 | 9.60e+02 | 1.63 | 1.22 | 1.06 | 1.00 | 0.93 | 0.90 | 0.89 | 0.76 |
| | | | 2048 | 1.92e+03 | 1.57 | 1.18 | 1.02 | 0.96 | 0.93 | 0.87 | 0.89 | 0.76 |
| | | | 4096 | 3.78e+03 | 1.48 | 1.11 | 0.98 | 0.91 | 0.88 | 0.84 | 0.84 | 0.74 |
| | | | 8192 | 7.54e+03 | 1.42 | 1.09 | 0.96 | 0.90 | 0.87 | 0.83 | 0.84 | 0.73 |
| | | Qwen2.5-7B | 64 | 1.49e+04 | 1.40 | 1.08 | 0.96 | 0.90 | 0.87 | 0.82 | 0.83 | 0.72 |
| | | | 128 | 3.50e+02 | 1.49 | 1.16 | 1.03 | 0.98 | 0.93 | 0.93 | 0.92 | 0.81 |
| | | | 256 | 4.03e+02 | 1.44 | 1.13 | 1.02 | 0.97 | 0.93 | 0.93 | 0.92 | 0.80 |
| | | | 512 | 6.18e+02 | 1.41 | 1.09 | 0.99 | 0.93 | 0.90 | 0.87 | 0.85 | 0.77 |
| | | | 1024 | 1.15e+03 | 1.35 | 1.04 | 0.93 | 0.88 | 0.83 | 0.82 | 0.82 | 0.68 |
| | | | 2048 | 2.19e+03 | 1.25 | 0.98 | 0.88 | 0.80 | 0.77 | 0.75 | 0.75 | 0.63 |
| | | | 4096 | 4.39e+03 | 1.35 | 1.01 | 0.90 | 0.82 | 0.78 | 0.77 | 0.77 | 0.65 |
| | | | 8192 | 8.62e+03 | 1.36 | 1.02 | 0.90 | 0.84 | 0.80 | 0.78 | 0.77 | 0.63 |
| | | Qwen2.5-14B | 64 | 3.38e+04 | 1.33 | 1.01 | 0.90 | 0.84 | 0.80 | 0.78 | 0.77 | 0.61 |
| | | | 128 | 4.40e+02 | 1.66 | 1.30 | 1.18 | 1.13 | 1.07 | 1.04 | 1.04 | 0.91 |
| | | | 256 | 4.72e+02 | 1.52 | 1.22 | 1.12 | 1.07 | 1.02 | 0.99 | 0.98 | 0.87 |
| | | | 512 | 7.12e+02 | 1.40 | 1.08 | 0.96 | 0.92 | 0.88 | 0.86 | 0.84 | 0.75 |
| | | | 1024 | 1.30e+03 | 1.39 | 1.04 | 0.92 | 0.88 | 0.83 | 0.79 | 0.79 | 0.69 |
| | | | 2048 | 2.51e+03 | 1.36 | 1.01 | 0.88 | 0.83 | 0.79 | 0.75 | 0.77 | 0.67 |
| | | | 4096 | 4.99e+03 | 1.41 | 1.05 | 0.93 | 0.87 | 0.83 | 0.79 | 0.81 | 0.70 |
| | | | 8192 | 9.88e+03 | 1.36 | 1.03 | 0.92 | 0.86 | 0.82 | 0.79 | 0.79 | 0.69 |
| | | RTX4090 | 64 | 1.98e+04 | 1.36 | 1.02 | 0.91 | 0.86 | 0.81 | 0.78 | 0.79 | 0.69 |
| | | | 128 | 3.96e+04 | 1.35 | 1.02 | 0.91 | 0.84 | 0.81 | 0.78 | 0.79 | 0.68 |
| | | | 256 | 9.71e+01 | 1.27 | 1.04 | 0.99 | 0.93 | 0.97 | 1.01 | 0.89 | 0.86 |
| | | | 512 | 1.31e+02 | 1.28 | 1.09 | 1.06 | 0.97 | 0.99 | 0.95 | 0.98 | 0.71 |
| | | | 1024 | 2.22e+02 | 1.64 | 1.33 | 1.20 | 1.14 | 0.99 | 1.07 | 0.97 | 0.85 |
| | | | 2048 | 4.12e+02 | 1.73 | 1.37 | 1.24 | 1.12 | 1.02 | 1.04 | 1.05 | 0.85 |
| | | | 4096 | 7.97e+02 | 1.86 | 1.41 | 0.99 | 1.18 | 1.10 | 1.07 | 0.84 | 0.91 |
| | | | 8192 | 1.54e+03 | 1.85 | 1.20 | 1.28 | 1.18 | 1.11 | 1.07 | 1.10 | 0.95 |
| | | Llama3.2-3B | 64 | 3.09e+03 | 1.88 | 1.44 | 1.30 | 1.20 | 1.11 | 1.09 | 1.10 | 0.95 |
| | | | 128 | 6.19e+03 | 1.89 | 1.45 | 1.30 | 1.20 | 1.14 | 1.12 | 1.10 | 0.94 |
| | | | 256 | 1.23e+04 | 1.88 | 1.43 | 1.29 | 1.19 | 1.16 | 1.12 | 1.12 | 0.96 |
| | | | 512 | 2.37e+02 | 2.09 | 1.73 | 1.21 | 1.09 | 1.05 | 1.00 | 0.99 | 0.89 |
| | | | 1024 | 4.12e+02 | 1.81 | 1.32 | 1.07 | 1.04 | 0.97 | 0.96 | 0.76 | 0.77 |
| | | | 2048 | 3.78e+02 | 1.83 | 1.36 | 1.22 | 1.18 | 1.12 | 0.98 | 0.94 | 0.90 |
| | | | 4096 | 6.86e+02 | 1.77 | 1.33 | 1.20 | 1.14 | 1.05 | 1.04 | 1.05 | 0.85 |
| | | | 8192 | 1.29e+03 | 1.80 | 1.34 | 1.19 | 1.14 | 1.02 | 1.06 | 1.03 | 0.91 |
| | | Qwen2.5-7B | 64 | 2.53e+03 | 1.86 | 1.36 | 1.20 | 1.14 | 1.09 | 1.07 | 1.06 | 0.93 |
| | | | 128 | 5.11e+03 | 1.92 | 1.43 | 1.27 | 1.20 | 1.15 | 1.11 | 1.11 | 0.97 |
| | | | 256 | 1.02e+04 | 1.95 | 1.43 | 1.26 | 1.20 | 1.14 | 1.12 | 1.10 | 0.94 |
| | | | 512 | 2.03e+04 | 1.94 | 1.42 | 1.25 | 1.18 | 1.13 | 1.10 | 1.09 | 0.96 |
| | | | 1024 | 5.24e+02 | 1.87 | 1.32 | 1.15 | 1.09 | 1.05 | 0.93 | 0.84 | 0.89 |
| | | | 2048 | 5.66e+02 | 1.54 | 1.15 | 0.98 | 0.92 | 0.85 | 0.78 | 0.80 | 0.65 |

■ $\geq 1.5 \times$ ■ $\geq 1.0 \times$

Appendix B Model Kernel Bf16 (Page 2/4)

| GPU | Model | M | cuBLASLt Latency (μs) | cuSPARSELt Speedup Ratio | | | | | | | |
|-------------|-------------|-------|--------------------------|--------------------------|------|------|------|------|------|------|------|
| | | | | 2:4 | 2:6 | 2:8 | 2:10 | 2:12 | 2:14 | 2:16 | 2:∞ |
| A100 | Qwen2.5-14B | 256 | 8.82e+02 | 1.41 | 1.08 | 0.98 | 0.90 | 0.66 | 0.81 | 0.81 | 0.70 |
| | | 512 | 1.60e+03 | 1.40 | 1.08 | 0.98 | 0.89 | 0.86 | 0.82 | 0.69 | 0.71 |
| | | 1024 | 3.10e+03 | 1.39 | 1.06 | 0.96 | 0.88 | 0.86 | 0.82 | 0.80 | 0.70 |
| | | 2048 | 5.98e+03 | 1.37 | 1.05 | 0.95 | 0.88 | 0.84 | 0.81 | 0.80 | 0.70 |
| | | 4096 | 1.18e+04 | 1.38 | 1.05 | 0.95 | 0.88 | 0.85 | 0.82 | 0.81 | 0.69 |
| | | 8192 | 2.36e+04 | 1.43 | 1.07 | 0.98 | 0.90 | 0.87 | 0.84 | 0.83 | 0.71 |
| | | 16384 | 4.72e+04 | 1.42 | 1.08 | 0.97 | 0.90 | 0.86 | 0.84 | 0.82 | 0.71 |
| | | 64 | 5.69e+02 | 1.87 | 1.32 | 1.16 | 1.10 | 0.85 | 1.01 | 0.99 | 0.75 |
| | | 128 | 6.45e+02 | 1.71 | 1.29 | 1.15 | 1.09 | 1.03 | 0.92 | 0.95 | 0.84 |
| | | 256 | 1.03e+03 | 1.57 | 1.23 | 1.03 | 0.99 | 1.00 | 0.96 | 0.91 | 0.82 |
| H100 | Llama3.2-1B | 512 | 1.91e+03 | 1.53 | 1.20 | 1.06 | 0.99 | 0.97 | 0.94 | 0.90 | 0.79 |
| | | 1024 | 3.71e+03 | 1.55 | 1.20 | 1.05 | 0.99 | 0.97 | 0.94 | 0.89 | 0.80 |
| | | 2048 | 7.45e+03 | 1.60 | 1.24 | 1.09 | 1.02 | 1.01 | 0.97 | 0.93 | 0.83 |
| | | 4096 | 1.49e+04 | 1.69 | 1.31 | 1.14 | 1.07 | 1.03 | 1.01 | 0.97 | 0.85 |
| | | 8192 | 3.00e+04 | 1.70 | 1.31 | 1.14 | 1.08 | 1.06 | 1.03 | 0.97 | 0.87 |
| | | 16384 | 6.01e+04 | 1.71 | 1.31 | 1.15 | 1.07 | 1.05 | 1.02 | 0.97 | 0.86 |
| | | 64 | 7.57e+01 | 0.74 | 0.61 | 0.57 | 0.53 | 0.51 | 0.48 | 0.50 | 0.42 |
| | | 128 | 8.10e+01 | 0.77 | 0.65 | 0.60 | 0.56 | 0.53 | 0.50 | 0.53 | 0.46 |
| | | 256 | 9.37e+01 | 0.87 | 0.72 | 0.66 | 0.61 | 0.59 | 0.57 | 0.58 | 0.52 |
| | | 512 | 1.43e+02 | 1.09 | 0.83 | 0.76 | 0.71 | 0.69 | 0.65 | 0.67 | 0.60 |
| L1 | Llama3.2-3B | 1024 | 2.63e+02 | 1.27 | 0.98 | 0.89 | 0.83 | 0.80 | 0.75 | 0.77 | 0.62 |
| | | 2048 | 5.40e+02 | 1.41 | 1.09 | 0.98 | 0.90 | 0.87 | 0.76 | 0.84 | 0.65 |
| | | 4096 | 1.05e+03 | 1.43 | 1.09 | 0.98 | 0.91 | 0.85 | 0.82 | 0.84 | 0.64 |
| | | 8192 | 2.16e+03 | 1.52 | 1.11 | 1.00 | 0.91 | 0.87 | 0.84 | 0.83 | 0.70 |
| | | 16384 | 4.64e+03 | 1.59 | 1.17 | 1.06 | 0.97 | 0.92 | 0.87 | 0.90 | 0.70 |
| | | 64 | 1.22e+02 | 0.93 | 0.72 | 0.66 | 0.63 | 0.61 | 0.59 | 0.58 | 0.53 |
| | | 128 | 1.25e+02 | 0.93 | 0.73 | 0.66 | 0.64 | 0.63 | 0.60 | 0.60 | 0.55 |
| | | 256 | 1.42e+02 | 1.03 | 0.81 | 0.73 | 0.70 | 0.69 | 0.66 | 0.66 | 0.59 |
| | | 512 | 2.25e+02 | 1.21 | 0.95 | 0.84 | 0.81 | 0.79 | 0.74 | 0.73 | 0.66 |
| | | 1024 | 4.46e+02 | 1.48 | 1.12 | 0.99 | 0.94 | 0.91 | 0.86 | 0.85 | 0.75 |
| Owen2.5-7B | Qwen2.5-7B | 2048 | 8.93e+02 | 1.56 | 1.18 | 1.04 | 0.97 | 0.95 | 0.88 | 0.88 | 0.78 |
| | | 4096 | 1.80e+03 | 1.60 | 1.15 | 1.01 | 0.94 | 0.91 | 0.87 | 0.86 | 0.73 |
| | | 8192 | 3.77e+03 | 1.61 | 1.17 | 1.03 | 0.99 | 0.93 | 0.89 | 0.86 | 0.75 |
| | | 16384 | 7.48e+03 | 1.53 | 1.13 | 1.01 | 0.91 | 0.91 | 0.84 | 0.83 | 0.73 |
| | | 64 | 2.69e+02 | 1.03 | 0.86 | 0.80 | 0.76 | 0.73 | 0.73 | 0.71 | 0.62 |
| | | 128 | 2.76e+02 | 1.13 | 0.90 | 0.85 | 0.80 | 0.77 | 0.76 | 0.74 | 0.66 |
| | | 256 | 3.05e+02 | 1.25 | 0.98 | 0.90 | 0.85 | 0.81 | 0.80 | 0.78 | 0.69 |
| | | 512 | 5.39e+02 | 1.43 | 1.09 | 0.97 | 0.93 | 0.87 | 0.86 | 0.84 | 0.74 |
| | | 1024 | 1.14e+03 | 1.74 | 1.27 | 1.15 | 1.07 | 1.02 | 1.01 | 0.98 | 0.76 |
| | | 2048 | 2.02e+03 | 1.54 | 1.11 | 0.97 | 0.91 | 0.83 | 0.81 | 0.80 | 0.71 |
| Qwen2.5-14B | Qwen2.5-14B | 4096 | 4.76e+03 | 1.65 | 1.26 | 1.11 | 1.03 | 0.92 | 0.93 | 0.92 | 0.78 |
| | | 8192 | 9.70e+03 | 1.58 | 1.24 | 1.09 | 0.99 | 0.99 | 0.95 | 0.92 | 0.80 |
| | | 16384 | 1.95e+04 | 1.71 | 1.20 | 1.10 | 0.99 | 0.90 | 0.89 | 0.88 | 0.76 |
| | | 64 | 3.19e+02 | 1.20 | 0.92 | 0.83 | 0.78 | | | | |
| | | 128 | 3.28e+02 | 1.27 | 0.96 | 0.88 | 0.83 | | | | |
| | | 256 | 3.68e+02 | 1.37 | 1.03 | 0.94 | 0.90 | | | | |
| | | 512 | 5.96e+02 | 1.41 | 1.08 | 0.97 | 0.91 | | | | |
| | | 1024 | 1.29e+03 | 1.64 | 1.24 | 1.10 | 1.02 | | | | |
| | | 2048 | 2.54e+03 | 1.62 | 1.18 | 1.08 | 0.99 | | | | |
| | | 4096 | 5.50e+03 | 1.66 | 1.25 | 1.06 | 1.03 | | | | |
| B200 | Llama3.2-1B | 8192 | 1.14e+04 | 1.66 | 1.21 | 1.09 | 1.04 | | | | |
| | | 16384 | 2.32e+04 | 1.60 | 1.18 | 1.13 | 1.00 | | | | |
| | | 64 | 4.46e+01 | 1.14 | 1.03 | 1.03 | 0.93 | | | | |
| | | 128 | 4.85e+01 | 1.18 | 1.09 | 1.07 | 0.96 | | | | |
| | | 256 | 6.07e+01 | 1.38 | 1.25 | 1.13 | 1.07 | | | | |
| | | 512 | 6.49e+01 | 1.24 | 1.05 | 0.98 | 0.90 | | | | |
| | | 1024 | 1.01e+02 | 1.36 | 1.09 | 1.03 | 0.92 | | | | |
| | | 2048 | 1.71e+02 | 1.41 | 1.13 | 1.03 | 0.94 | | | | |
| | | 4096 | 3.21e+02 | 1.50 | 1.16 | 1.07 | 0.95 | | | | |
| | | 8192 | 6.49e+02 | 1.59 | 1.18 | 1.08 | 0.96 | | | | |
| L1 | Llama3.2-3B | 16384 | 1.34e+03 | 1.67 | 1.29 | 1.20 | 1.04 | | | | |
| | | 64 | 5.68e+01 | 1.37 | 1.18 | 1.15 | 1.10 | | | | |
| | | 128 | 5.80e+01 | 1.29 | 1.12 | 1.04 | 0.97 | | | | |
| | | 256 | 6.21e+01 | 1.20 | 1.00 | 0.97 | 0.87 | | | | |
| | | 512 | 9.06e+01 | 1.33 | 1.07 | 0.98 | 0.92 | | | | |

Appendix B Model Kernel Bf16 (Page 3/4)

| GPU | Model | M | cuBLASLt Latency (μs) | cuSPARSELt Speedup Ratio | | | | | | | |
|-------------|-------------|-------|--------------------------|--------------------------|------|------|------|------|------|------|------|
| | | | | 2:4 | 2:6 | 2:8 | 2:10 | 2:12 | 2:14 | 2:16 | 2:∞ |
| Qwen2.5-7B | Qwen2.5-7B | 1024 | 1.47e+02 | 1.33 | 1.07 | 0.97 | 0.90 | | | | |
| | | 2048 | 2.67e+02 | 1.43 | 1.11 | 1.00 | 0.92 | | | | |
| | | 4096 | 5.22e+02 | 1.59 | 1.20 | 1.08 | 1.00 | | | | |
| | | 8192 | 1.07e+03 | 1.67 | 1.25 | 1.11 | 1.05 | | | | |
| | | 16384 | 2.13e+03 | 1.68 | 1.30 | 1.14 | 1.04 | | | | |
| | | 64 | 1.02e+02 | 1.48 | 1.19 | 1.08 | 1.02 | | | | |
| | | 128 | 1.05e+02 | 1.36 | 1.12 | 1.01 | 0.97 | | | | |
| | | 256 | 1.22e+02 | 1.40 | 1.07 | 1.01 | 0.97 | | | | |
| | | 512 | 1.82e+02 | 1.41 | 1.10 | 1.02 | 0.96 | | | | |
| | | 1024 | 3.18e+02 | 1.53 | 1.18 | 1.09 | 1.02 | | | | |
| Qwen2.5-14B | Qwen2.5-14B | 2048 | 6.37e+02 | 1.66 | 1.26 | 1.14 | 1.08 | | | | |
| | | 4096 | 1.27e+03 | 1.70 | 1.27 | 1.17 | 1.10 | | | | |
| | | 8192 | 2.64e+03 | 1.81 | 1.38 | 1.24 | 1.14 | | | | |
| | | 16384 | 5.27e+03 | 1.81 | 1.33 | 1.20 | 1.10 | | | | |
| | | 64 | 1.13e+02 | 1.57 | 1.20 | 1.10 | 1.03 | | | | |
| | | 128 | 1.16e+02 | 1.43 | 1.10 | 1.05 | 0.97 | | | | |
| | | 256 | 1.36e+02 | 1.32 | 1.05 | 0.98 | 0.92 | | | | |
| | | 512 | 2.02e+02 | 1.35 | 1.06 | 0.98 | 0.92 | | | | |
| | | 1024 | 4.00e+02 | 1.60 | 1.21 | 1.10 | 1.04 | | | | |
| | | 2048 | 7.21e+02 | 1.59 | 1.18 | 1.08 | 1.00 | | | | |
| RTX5080 | Llama3.2-1B | 4096 | 1.44e+03 | 1.65 | 1.23 | 1.09 | 1.01 | | | | |
| | | 8192 | 3.04e+03 | 1.80 | 1.31 | 1.20 | 1.08 | | | | |
| | | 16384 | 6.14e+03 | 1.77 | 1.28 | 1.13 | 1.06 | | | | |
| | | 64 | 1.13e+02 | 1.31 | 1.05 | 0.99 | 0.97 | 0.94 | 0.95 | 0.90 | 0.67 |
| | | 128 | 1.98e+02 | 1.64 | 1.30 | 1.18 | 1.14 | 1.07 | 1.05 | 1.05 | 0.91 |
| | | 256 | 3.60e+02 | 1.79 | 1.40 | 1.26 | 1.19 | 1.14 | 1.13 | 1.11 | 0.97 |
| | | 512 | 6.38e+02 | 1.81 | 1.40 | 1.27 | 1.18 | 1.14 | 1.12 | 1.10 | 0.98 |
| | | 1024 | 1.23e+03 | 1.91 | 1.47 | 1.33 | 1.24 | 1.20 | 1.17 | 1.15 | 1.02 |
| | | 2048 | 2.31e+03 | 1.87 | 1.44 | 1.29 | 1.20 | 1.16 | 1.13 | 1.12 | 0.98 |
| | | 4096 | 4.41e+03 | 1.85 | 1.41 | 1.28 | 1.18 | 1.14 | 1.11 | 1.09 | 0.96 |
| BitNet-2B | BitNet-2B | 8192 | 8.63e+03 | 1.87 | 1.42 | 1.28 | 1.18 | 1.13 | 1.11 | 1.09 | 0.95 |
| | | 16384 | 1.70e+04 | 1.87 | 1.43 | 1.27 | 1.17 | 1.13 | 1.11 | 1.09 | 0.95 |
| | | 64 | 2.47e+02 | 2.96 | 2.46 | 2.28 | 2.19 | 2.03 | 1.88 | 1.82 | 1.43 |
| | | 128 | 2.91e+02 | 2.39 | 1.87 | 1.72 | 1.62 | 1.55 | 1.46 | 1.43 | 1.22 |
| | | 256 | 3.60e+02 | 1.78 | 1.39 | 1.25 | 1.17 | 1.12 | 1.09 | 1.06 | 0.94 |
| | | 512 | 7.01e+02 | 1.91 | 1.45 | 1.31 | 1.23 | 1.17 | 1.14 | 1.13 | 0.98 |
| | | 1024 | 1.29e+03 | 1.89 | 1.43 | 1.29 | 1.22 | 1.16 | 1.12 | 1.10 | 0.96 |
| | | 2048 | 2.49e+03 | 1.87 | 1.43 | 1.29 | 1.21 | 1.16 | 1.12 | 1.11 | 0.97 |
| | | 4096 | 4.95e+03 | 1.89 | 1.43 | 1.29 | 1.20 | 1.14 | 1.10 | 1.09 | 0.95 |
| | | 8192 | 9.82e+03 | 1.89 | 1.43 | 1.28 | 1.20 | 1.14 | 1.10 | 1.09 | 0.95 |
| Llama3.2-3B | Llama3.2-3B | 16384 | 1.93e+04 | 1.89 | 1.42 | 1.28 | 1.19 | 1.13 | 1.09 | 1.08 | 0.94 |
| | | 64 | 3.80e+02 | 3.24 | 2.23 | 2.04 | 1.89 | 1.85 | 1.80 | 1.80 | 1.61 |
| | | 128 | 4.26e+02 | 2.55 | 1.86 | 1.72 | 1.60 | 1.54 | 1.50 | 1.48 | 1.33 |
| | | 256 | 5.61e+02 | 1.88 | 1.40 | 1.27 | 1.20 | 1.14 | 1.09 | 1.05 | 0.92 |
| | | 512 | 1.05e+03 | 1.90 | 1.47 | 1.31 | 1.23 | 1.17 | 1.12 | 1.08 | 0.90 |
| | | 1024 | 1.91e+03 | 1.90 | 1.46 | 1.30 | 1.20 | 1.14 | 1.09 | 1.05 | 0.88 |
| | | 2048 | 3.63e+03 | 1.90 | 1.43 | 1.27 | 1.18 | 1.13 | 1.07 | 1.04 | 0.87 |
| | | 4096 | 7.19e+03 | 1.90 | 1.41 | 1.26 | 1.17 | 1.12 | 1.06 | 1.04 | 0.87 |
| | | 8192 | 1.40e+04 | 1.89 | 1.40 | 1.25 | 1.14 | 1.09 | 1.04 | 1.02 | 0.87 |
| | | 16384 | 2.78e+04 | 1.89 | 1.40 | 1.23 | 1.12 | 1.08 | 1.02 | 1.01 | 0.87 |
| Qwen2.5-7B | Qwen2.5-7B | 64 | 8.98e+02 | 2.92 | 2.13 | 1.91 | 1.79 | 1.69 | 1.68 | 1.63 | 1.42 |
| | | 128 | 6.30e+02 | 1.78 | 1.31 | 1.17 | 1.10 | 1.05 | 1.02 | 1.00 | 0.87 |
| | | 256 | 1.13e+03 | 1.66 | 1.27 | 1.11 | 1.03 | 0.96 | 0.94 | 0.94 | 0.80 |
| | | 512 | 2.16e+03 | 1.69 | 1.28 | 1.13 | 1.05 | 0.98 | 0.97 | 0.94 | 0.81 |
| | | 1024 | 4.25e+03 | 1.72 | 1.29 | 1.15 | 1.07 | 1.01 | 0.99 | 0.96 | 0.83 |
| | | 2048 | 8.18e+03 | 1.69 | 1.27 | 1.13 | 1.06 | 0.99 | 0.97 | 0.95 | 0.81 |
| | | 4096 | 1.62e+04 | 1.70 | 1.29 | 1.15 | 1.07 | 1.00 | 0.99 | 0.95 | 0.82 |
| | | 8192 | 3.21e+04 | 1.69 | 1.28 | 1.14 | 1.07 | 1.00 | 0.98 | 0.96 | 0.82 |
| | | 16384 | 6.39e+04 | 1.69 | 1.28 | 1.14 | 1.07 | 1.01 | 0.98 | 0.96 | 0.81 |
| | | 64 | 1.07e+03 | 3.08 | 2.30 | 2.04 | 1.94 | 1.86 | 1.76 | 1.72 | 1.43 |
| Qwen2.5-14B | Qwen2.5-14B | 128 | 1.03e+03 | 2.45 | 1.80 | 1.64 | 1.54 | 1.49 | 1.41 | 1.40 | 1.20 |
| | | 256 | 1.38e+03 | 1.79 | 1.31 | 1.19 | 1.10 | 1.06 | 1.02 | 1.01 | 0.88 |
| | | 512 | 2.51e+03 | 1.73 | 1.26 | 1.13 | 1.04 | 1.02 | 0.98 | 0.96 | 0.84 |
| | | 1024 | 4.88e+03 | 1.73 | 1.25 | 1.13 | 1.05 | 1.01 | 0.97 | 0.96 | 0.84 |
| | | 2048 | 9.66e+03 | 1.74 | 1.24 | 1.12 | 1.04 | 1.00 | 0.96 | 0.95 | 0.83 |

Appendix B Model Kernel Bf16 (Page 4/4)

| GPU | Model | M | cuBLASLt Latency (μs) | cuSPARSELt Speedup Ratio | | | | | | | |
|-------------|-------------|-------|--------------------------|--------------------------|------|------|------|------|------|------|------|
| | | | | 2:4 | 2:6 | 2:8 | 2:10 | 2:12 | 2:14 | 2:16 | 2:∞ |
| GB10 | Llama3.2-1B | 4096 | 1.92e+04 | 1.73 | 1.24 | 1.11 | 1.03 | 1.00 | 0.96 | 0.95 | 0.83 |
| | | 8192 | 3.78e+04 | 1.72 | 1.23 | 1.10 | 1.02 | 0.99 | 0.95 | 0.94 | 0.83 |
| | | 16384 | 7.49e+04 | 1.71 | 1.22 | 1.10 | 1.01 | 0.99 | 0.95 | 0.94 | 0.82 |
| | | 64 | 5.78e+02 | 2.66 | 1.75 | 1.45 | 1.08 | 0.97 | 1.09 | 1.12 | 0.96 |
| | | 128 | 7.01e+02 | 2.69 | 1.84 | 1.47 | 1.11 | 1.04 | 1.17 | 1.19 | 1.05 |
| | | 256 | 7.17e+02 | 1.45 | 1.01 | 0.81 | 0.64 | 0.60 | 0.68 | 0.70 | 0.59 |
| | | 512 | 9.84e+02 | 1.06 | 0.67 | 0.64 | 0.47 | 0.42 | 0.48 | 0.50 | 0.43 |
| | | 1024 | 1.61e+03 | 0.81 | 0.59 | 0.51 | 0.38 | 0.37 | 0.41 | 0.43 | 0.37 |
| | | 2048 | 2.97e+03 | 0.79 | 0.52 | 0.50 | 0.36 | 0.33 | 0.39 | 0.40 | 0.34 |
| | | 4096 | 5.90e+03 | 0.78 | 0.54 | 0.48 | 0.36 | 0.33 | 0.39 | 0.42 | 0.34 |
| | | 8192 | 1.12e+04 | 0.72 | 0.51 | 0.45 | 0.34 | 0.31 | 0.38 | 0.38 | 0.33 |
| | | 16384 | 2.20e+04 | 0.75 | 0.51 | 0.44 | 0.34 | 0.33 | 0.37 | 0.39 | 0.32 |
| Llama3.2-3B | Llama3.2-3B | 64 | 9.97e+02 | 2.50 | 1.59 | 1.34 | 1.27 | 1.12 | 0.90 | 1.10 | 0.91 |
| | | 128 | 1.17e+03 | 2.34 | 1.57 | 1.40 | 1.28 | 1.12 | 0.94 | 1.16 | 1.02 |
| | | 256 | 1.20e+03 | 1.37 | 0.91 | 0.80 | 0.72 | 0.68 | 0.57 | 0.70 | 0.60 |
| | | 512 | 1.44e+03 | 0.85 | 0.56 | 0.49 | 0.46 | 0.41 | 0.33 | 0.43 | 0.35 |
| | | 1024 | 2.56e+03 | 0.74 | 0.51 | 0.45 | 0.41 | 0.37 | 0.30 | 0.38 | 0.33 |
| | | 2048 | 4.99e+03 | 0.72 | 0.50 | 0.45 | 0.42 | 0.38 | 0.30 | 0.38 | 0.32 |
| | | 4096 | 9.88e+03 | 0.73 | 0.51 | 0.45 | 0.42 | 0.38 | 0.30 | 0.38 | 0.32 |
| | | 8192 | 1.86e+04 | 0.71 | 0.49 | 0.42 | 0.40 | 0.36 | 0.28 | 0.36 | 0.30 |
| | | 16384 | 3.68e+04 | 0.70 | 0.48 | 0.42 | 0.39 | 0.36 | 0.28 | 0.36 | 0.30 |
| | | 64 | 2.38e+03 | 1.95 | 1.41 | 1.25 | 1.15 | | | | |
| Qwen2.5-7B | Qwen2.5-7B | 128 | 2.59e+03 | 1.92 | 1.40 | 1.24 | 1.15 | | | | |
| | | 256 | 2.87e+03 | 1.18 | 0.85 | 0.76 | 0.69 | | | | |
| | | 512 | 3.75e+03 | 0.79 | 0.58 | 0.51 | 0.47 | | | | |
| | | 1024 | 6.16e+03 | 0.66 | 0.48 | 0.43 | 0.40 | | | | |
| | | 2048 | 1.19e+04 | 0.66 | 0.48 | 0.42 | 0.39 | | | | |
| | | 4096 | 2.26e+04 | 0.63 | 0.46 | 0.41 | 0.37 | | | | |
| | | 8192 | 4.29e+04 | 0.59 | 0.44 | 0.38 | 0.36 | | | | |
| | | 16384 | 8.46e+04 | 0.59 | 0.43 | 0.38 | 0.35 | | | | |
| | | 64 | 3.02e+03 | 1.99 | 1.46 | 1.29 | 1.20 | | | | |
| | | 128 | 3.14e+03 | 1.89 | 1.33 | 1.23 | 1.12 | | | | |
| Qwen2.5-14B | Qwen2.5-14B | 256 | 3.40e+03 | 1.09 | 0.77 | 0.69 | 0.66 | | | | |
| | | 512 | 4.50e+03 | 0.70 | 0.52 | 0.47 | 0.44 | | | | |
| | | 1024 | 7.64e+03 | 0.61 | 0.45 | 0.40 | 0.37 | | | | |
| | | 2048 | 1.55e+04 | 0.64 | 0.46 | 0.41 | 0.38 | | | | |
| | | 4096 | 2.88e+04 | 0.60 | 0.43 | 0.39 | 0.36 | | | | |
| | | 8192 | 5.55e+04 | 0.58 | 0.42 | 0.37 | 0.35 | | | | |
| | | 16384 | 1.09e+05 | 0.56 | 0.41 | 0.36 | 0.34 | | | | |